

Assignment is below at the end

- <https://scikit-learn.org/stable/modules/tree.html>
- <https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html>
- https://scikit-learn.org/stable/modules/generated/sklearn.tree.plot_tree.html

```
In [325... import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
plt.rcParams['figure.figsize'] = (20, 6)
plt.rcParams['font.size'] = 14
import pandas as pd
```

```
In [326... df = pd.read_csv('adult.data', index_col=False)
```

```
In [327... golden = pd.read_csv('adult.test', index_col=False)
```

```
In [328... golden.head()
```

Out[328]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male
4	18	?	103497	Some-college	10	Never-married	?	Own-child	White	Female

```
In [329... df.head()
```

Out[329]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female

In [330... `df.columns`

Out[330]: Index(['age', 'workclass', 'fnlwgt', 'education', 'education-num', 'marital-status', 'occupation', 'relationship', 'race', 'sex', 'capital-gain', 'capital-loss', 'hours-per-week', 'native-country', 'salary'], dtype='object')

In [331... `from sklearn import preprocessing`

```
In [332... # Columns we want to transform
transform_columns = ['sex']

#Columns we can't use because non-numerical
non_num_columns = ['workclass', 'education', 'marital-status',
                   'occupation', 'relationship', 'race', 'sex',
                   'native-country']
```

First let's try using `pandas.get_dummies()` to transform columns

```
In [333... dummies = pd.get_dummies(df[transform_columns])
dummies
```

Out[333]:

	sex_Female	sex_Male
0	0	1
1	0	1
2	0	1
3	0	1
4	1	0
...
32556	1	0
32557	0	1
32558	1	0
32559	0	1
32560	1	0

32561 rows × 2 columns

In [334... `dummies.shape`

Out[334]: (32561, 2)

sklearn has a similar process for OneHot Encoding features

In [335... `onehot = preprocessing.OneHotEncoder(handle_unknown = "infrequent_if_exist", sparse=False)`
`onehot.fit(df[transform_columns])`

C:\Users\jorda\anaconda3\lib\site-packages\sklearn\preprocessing_encoders.py:828: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.
 warnings.warn(

Out[335]:

▼ OneHotEncoder

OneHotEncoder(handle_unknown='infrequent_if_exist', sparse=False, sparse_output=False)

In [336... `onehot.categories_`

Out[336]: [array([' Female', ' Male'], dtype=object)]

In [337... `sex = onehot.transform(df[transform_columns])`
`sex`

```
Out[337]: array([[0., 1.],
          [0., 1.],
          [0., 1.],
          ...,
          [1., 0.],
          [0., 1.],
          [1., 0.]])
```

```
In [338... sex.shape
```

```
Out[338]: (32561, 2)
```

In addition to OneHot encoding there is Ordinal Encoding

```
In [339... enc = preprocessing.OrdinalEncoder()
enc.fit(df[["salary"]])
salary = enc.transform(df[["salary"]])
salary
```

```
Out[339]: array([[0.],
          [0.],
          [0.],
          ...,
          [0.],
          [0.],
          [1.]])
```

```
In [340... enc.categories_[0]
```

```
Out[340]: array([' <=50K', ' >50K'], dtype=object)
```

```
In [341... x = df.copy()

# transformed = pd.get_dummies(df[transform_columns])

onehot = preprocessing.OneHotEncoder(handle_unknown="infrequent_if_exist", sparse=False)

enc = preprocessing.OrdinalEncoder()

enc.fit(df[["salary"]])

transformed = onehot.transform(df[transform_columns])
new_cols = list(onehot.categories_[0].flatten())
df_trans = pd.DataFrame(transformed, columns=new_cols)

x = pd.concat(
    [
        x.drop(non_num_columns, axis=1),
        df_trans
    ],
    axis=1,)
```

```
x["salary"] = enc.transform(df[["salary"]])
```

C:\Users\jorda\anaconda3\lib\site-packages\sklearn\preprocessing_encoders.py:828: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.

warnings.warn(

In [342...

x

Out[342]:

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week	salary	Female	Male
0	39	77516	13	2174	0	40	0.0	0.0	1.0
1	50	83311	13	0	0	13	0.0	0.0	1.0
2	38	215646	9	0	0	40	0.0	0.0	1.0
3	53	234721	7	0	0	40	0.0	0.0	1.0
4	28	338409	13	0	0	40	0.0	1.0	0.0
...
32556	27	257302	12	0	0	38	0.0	1.0	0.0
32557	40	154374	9	0	0	40	1.0	0.0	1.0
32558	58	151910	9	0	0	40	0.0	1.0	0.0
32559	22	201490	9	0	0	20	0.0	0.0	1.0
32560	52	287927	9	15024	0	40	1.0	1.0	0.0

32561 rows × 9 columns

In [343...

```
xt = golden.copy()

transformed = onehot.transform(xt[transform_columns])
new_cols = list(onehot.categories_[0].flatten())
df_transx = pd.DataFrame(transformed, columns=new_cols)

xt = pd.concat(
    [
        xt.drop(non_num_columns, axis=1),
        df_transx
    ],
    axis=1,
)

xt["salary"] = enc.fit_transform(golden[["salary"]])
```

In [344...

```
xt.salary.value_counts()
```

Out[344]:

```
0.0    12435
1.0     3846
Name: salary, dtype: int64
```

In [345...

```
enc.categories_
```

Out[345]:

```
[array([' <=50K.', ' >50K.'], dtype=object)]
```

```
In [346... from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
```

Choose the model of your preference: DecisionTree or RandomForest

```
In [347... model = RandomForestClassifier(criterion='entropy')
```

```
In [348... model = DecisionTreeClassifier(criterion='entropy', max_depth=None)
```

```
In [349... model.fit(x.drop(['fnlwgt', 'salary'], axis=1), x.salary)
```

```
Out[349]: DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy')
```

```
In [350... model.tree_.node_count
```

```
Out[350]: 8313
```

```
In [351... list(zip(x.drop(['fnlwgt', 'salary'], axis=1).columns, model.feature_importances_))
```

```
Out[351]: [('age', 0.32449288950363436),
('education-num', 0.16041765339065917),
('capital-gain', 0.22731243895177794),
('capital-loss', 0.07830610027050318),
('hours-per-week', 0.15395455662211974),
(' Female', 0.03396418514452267),
(' Male', 0.021552176116782933)]
```

```
In [352... list(zip(x.drop(['fnlwgt', 'salary'], axis=1).columns, model.feature_importances_))
```

```
Out[352]: [('age', 0.32449288950363436),
('education-num', 0.16041765339065917),
('capital-gain', 0.22731243895177794),
('capital-loss', 0.07830610027050318),
('hours-per-week', 0.15395455662211974),
(' Female', 0.03396418514452267),
(' Male', 0.021552176116782933)]
```

```
In [353... x.drop(['fnlwgt', 'salary'], axis=1).head()
```

```
Out[353]:
```

	age	education-num	capital-gain	capital-loss	hours-per-week	Female	Male
0	39	13	2174	0	40	0.0	1.0
1	50	13	0	0	13	0.0	1.0
2	38	9	0	0	40	0.0	1.0
3	53	7	0	0	40	0.0	1.0
4	28	13	0	0	40	1.0	0.0

```
In [354... set(x.columns) - set(xt.columns)
```

Out[354]: set()

In [355... list(x.drop('salary', axis=1).columns)

Out[355]:

```
['age',
 'fnlwgt',
 'education-num',
 'capital-gain',
 'capital-loss',
 'hours-per-week',
 ' Female',
 ' Male']
```

In [356... predictions = model.predict(xt.drop(['fnlwgt', 'salary'], axis=1))
 predictionxs = model.predict(x.drop(['fnlwgt', 'salary'], axis=1))

In [357... from sklearn.metrics import (
 accuracy_score,
 classification_report,
 confusion_matrix, auc, roc_curve
)

In [358... accuracy_score(xt.salary, predictions)

Out[358]: 0.8208341010994411

In [359... confusion_matrix(xt.salary, predictions)

Out[359]: array([[11460, 975],
 [1942, 1904]], dtype=int64)

In [360... print(classification_report(xt.salary, predictions))

	precision	recall	f1-score	support
0.0	0.86	0.92	0.89	12435
1.0	0.66	0.50	0.57	3846
accuracy			0.82	16281
macro avg	0.76	0.71	0.73	16281
weighted avg	0.81	0.82	0.81	16281

In [361... print(classification_report(xt.salary, predictions))

	precision	recall	f1-score	support
0.0	0.86	0.92	0.89	12435
1.0	0.66	0.50	0.57	3846
accuracy			0.82	16281
macro avg	0.76	0.71	0.73	16281
weighted avg	0.81	0.82	0.81	16281

In [362... accuracy_score(x.salary, predictionxs)

Out[362]: 0.8955806025613464

```
In [363... confusion_matrix(x.salary, predictionsx)
```

```
Out[363]: array([[24097,   623],
        [ 2777,  5064]], dtype=int64)
```

```
In [364... print(classification_report(x.salary, predictionsx))
```

	precision	recall	f1-score	support
0.0	0.90	0.97	0.93	24720
1.0	0.89	0.65	0.75	7841
accuracy			0.90	32561
macro avg	0.89	0.81	0.84	32561
weighted avg	0.90	0.90	0.89	32561

```
In [365... print(classification_report(x.salary, predictionsx))
```

	precision	recall	f1-score	support
0.0	0.90	0.97	0.93	24720
1.0	0.89	0.65	0.75	7841
accuracy			0.90	32561
macro avg	0.89	0.81	0.84	32561
weighted avg	0.90	0.90	0.89	32561

For the following use the above `adult` dataset.

1. Show the RandomForest outperforms the DecisionTree for a fixed `max_depth` by training using the train set and calculate precision, recall, f1, confusion matrix on golden-test set. Start with only numerical features/columns. (age, education-num, capital-gain, capital-loss, hours-per-week)

```
In [366... modela = RandomForestClassifier(criterion = 'entropy', max_depth = 5)
        modelb = DecisionTreeClassifier(criterion = 'entropy', max_depth = 5)
```

```
In [367... modela.fit(x.drop(['fnlwtg', 'salary'], axis = 1), x.salary)
        modelb.fit(x.drop(['fnlwtg', 'salary'], axis = 1), x.salary)
```


Out[367]:

```
DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy', max_depth=5)
```

In [368...

```
predictionsa = modela.predict(xt.drop(['fnlwgt', 'salary'], axis = 1))
predictionsax = modela.predict(x.drop(['fnlwgt', 'salary'], axis = 1))
predictionsb = modelb.predict(xt.drop(['fnlwgt', 'salary'], axis = 1))
predictionsbx = modelb.predict(x.drop(['fnlwgt', 'salary'], axis = 1))
```

In [369...

```
accuracy_score(xt.salary, predictionsa)
```

Out[369]:

```
0.8312142988759904
```

In [370...

```
accuracy_score(xt.salary, predictionsb)
```

Out[370]:

```
0.8201584669246361
```

In [371...

```
confusion_matrix(xt.salary, predictionsa)
```

Out[371]:

```
array([[12049,  386],
       [ 2362, 1484]], dtype=int64)
```

In [372...

```
confusion_matrix(xt.salary, predictionsb)
```

Out[372]:

```
array([[11458,  977],
       [ 1951, 1895]], dtype=int64)
```

In [373...

```
print(classification_report(xt.salary, predictionsa))
```

	precision	recall	f1-score	support
0.0	0.84	0.97	0.90	12435
1.0	0.79	0.39	0.52	3846
accuracy			0.83	16281
macro avg	0.81	0.68	0.71	16281
weighted avg	0.83	0.83	0.81	16281

In [374...

```
print(classification_report(xt.salary, predictionsb))
```

	precision	recall	f1-score	support
0.0	0.85	0.92	0.89	12435
1.0	0.66	0.49	0.56	3846
accuracy			0.82	16281
macro avg	0.76	0.71	0.73	16281
weighted avg	0.81	0.82	0.81	16281

2. Use a RandomForest or DecisionTree and the `adult` dataset, systematically add new columns, one by one, that are non-numerical

but converted using the feature-extraction techniques we learned. Using the golden-test set show [precision, recall, f1, confusion matrix] for each additional feature added.

```
In [375... ##Workclass
transform_columns1 = ['workclass']
non_num_columns1 = ['education', 'marital-status',
                    'occupation', 'relationship', 'race', 'sex',
                    'native-country']
```

```
In [397... data1 = df.copy()

onehot1 = preprocessing.OneHotEncoder(handle_unknown="infrequent_if_exist", sparse=False)

enc1 = preprocessing.OrdinalEncoder()

enc1.fit(df[["workclass"]])

transformed1 = onehot1.transform(df[transform_columns1])
new_cols1 = list(onehot1.categories_[0].flatten())
df_trans1 = pd.DataFrame(transformed1, columns = new_cols1)

data1 = pd.concat(
    [
        data1.drop(non_num_columns1, axis = 1),
        df_trans,
        df_trans1
    ],
    axis = 1,)

data1["workclass"] = enc1.fit_transform(df[["workclass"]])
data1["sex"] = enc1.fit_transform(df[["sex"]])
```

C:\Users\jorda\anaconda3\lib\site-packages\sklearn\preprocessing_encoders.py:828: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.

```
warnings.warn(
```

```
In [377... data1x = golden.copy()

transformed1 = onehot1.transform(golden[transform_columns1])
new_cols1 = list(onehot1.categories_[0].flatten())
df_trans1x = pd.DataFrame(transformed1, columns = new_cols1)

data1x = pd.concat(
    [
        data1x.drop(non_num_columns1, axis = 1),
        df_transx,
        df_trans1x
```

```

    ],
    axis = 1,)

data1x["workclass"] = enc1.fit_transform(golden[["workclass"]])
data1x["sex"] = enc1.fit_transform(golden[["sex"]])

```

```
In [378...] modelc = RandomForestClassifier(criterion = 'entropy', max_depth = 5)
```

```
In [379...] modelc.fit(data1.drop(['fnlwgt', 'salary'], axis = 1), data1.salary)
```

```
Out[379]: ▼ RandomForestClassifier
RandomForestClassifier(criterion='entropy', max_depth=5)
```

```
In [380...] predictionsc = modelc.predict(data1x.drop(['fnlwgt', 'salary'], axis = 1))
predictionscx = modelc.predict(data1.drop(['fnlwgt', 'salary'], axis = 1))
```

```
In [381...] confusion_matrix(data1.salary, predictionscx)
```

```
Out[381]: array([[24196,   524],
 [ 4879,  2962]], dtype=int64)
```

```
In [382...] print(classification_report(data1.salary, predictionscx))
```

	precision	recall	f1-score	support
<=50K	0.83	0.98	0.90	24720
>50K	0.85	0.38	0.52	7841
accuracy			0.83	32561
macro avg	0.84	0.68	0.71	32561
weighted avg	0.84	0.83	0.81	32561

```
In [435...] ##Marital Status
transform_columns2 = ['marital-status']
non_num_columns2 = ['education', 'occupation', 'relationship', 'race', 'sex', 'native-
```

```
In [426...] data2 = df.copy()

onehot2 = preprocessing.OneHotEncoder(handle_unknown="infrequent_if_exist", sparse=False)

enc2 = preprocessing.OrdinalEncoder()

enc2.fit(df[["marital-status"]])

transformed2 = onehot2.transform(df[transform_columns2])
new_cols2 = list(onehot2.categories_[0].flatten())
df_trans2 = pd.DataFrame(transformed2, columns = new_cols2)

data2 = pd.concat(
    [
        data2.drop(non_num_columns2, axis = 1),
        df_trans,
        df_trans1,
        df_trans2
    ]
)
```

```

],
axis = 1,)

data2["workclass"] = enc2.fit_transform(df[["workclass"]])
data2["marital-status"] = enc2.fit_transform(df[["marital-status"]])
data2["sex"] = enc2.fit_transform(df[["sex"]])

```

C:\Users\jorda\anaconda3\lib\site-packages\sklearn\preprocessing_encoders.py:828: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.

warnings.warn(

```

In [431... data2x = golden.copy()

transformed2 = onehot2.transform(golden[transform_columns2])
new_cols2 = list(onehot2.categories_[0].flatten())
df_trans2x = pd.DataFrame(transformed2, columns = new_cols2)

data2x = pd.concat(
    [
        data2x.drop(non_num_columns2, axis = 1),
        df_transx,
        df_trans1x,
        df_trans2x
    ],
    axis = 1,)

data2x["workclass"] = enc2.fit_transform(golden[["workclass"]])
data2x["marital-status"] = enc2.fit_transform(golden[["marital-status"]])
data2x["sex"] = enc2.fit_transform(golden[["sex"]])

```

```

In [440... modeld = RandomForestClassifier(criterion = 'entropy', max_depth = 5)

```

```

In [441... modeld.fit(data2.drop(['fnlwgt', 'salary'], axis = 1), data2.salary)

```

```

Out[441]: ▼ RandomForestClassifier
RandomForestClassifier(criterion='entropy', max_depth=5)

```

```

In [442... predictionsd = modeld.predict(data2x.drop(['fnlwgt', 'salary'], axis = 1))
predictionsdx = modeld.predict(data2.drop(['fnlwgt', 'salary'], axis = 1))

```

```

In [443... confusion_matrix(data2.salary, predictionsdx)

```

```

Out[443]: array([[23688, 1032],
       [ 3788, 4053]], dtype=int64)

```

```

In [444... print(classification_report(data2.salary, predictionsdx))

```

	precision	recall	f1-score	support
<=50K	0.86	0.96	0.91	24720
>50K	0.80	0.52	0.63	7841
accuracy			0.85	32561
macro avg	0.83	0.74	0.77	32561
weighted avg	0.85	0.85	0.84	32561

```

In [436... ##Education
transform_columns3 = ['education']
non_num_columns3 = ['occupation', 'relationship', 'race', 'sex', 'native-country']

In [437... data3 = df.copy()

onehot3 = preprocessing.OneHotEncoder(handle_unknown="infrequent_if_exist", sparse=False)

enc3 = preprocessing.OrdinalEncoder()

enc3.fit(df[["education"]])

transformed3 = onehot3.transform(df[transform_columns3])
new_cols3 = list(onehot3.categories_[0].flatten())
df_trans3 = pd.DataFrame(transformed3, columns = new_cols3)

data3 = pd.concat(
    [
        data3.drop(non_num_columns3, axis = 1),
        df_trans,
        df_trans1,
        df_trans2,
        df_trans3
    ],
    axis = 1,)

data3["workclass"] = enc3.fit_transform(df[["workclass"]])
data3["marital-status"] = enc3.fit_transform(df[["marital-status"]])
data3["education"] = enc3.fit_transform(df[["education"]])
data3["sex"] = enc3.fit_transform(df[["sex"]])

```

C:\Users\jorda\anaconda\lib\site-packages\sklearn\preprocessing_encoders.py:828: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.

warnings.warn(

```

In [439... data3x = golden.copy()

transformed3 = onehot3.transform(golden[transform_columns3])
new_cols3 = list(onehot3.categories_[0].flatten())
df_trans3x = pd.DataFrame(transformed3, columns = new_cols3)

data3x = pd.concat(
    [
        data3x.drop(non_num_columns3, axis = 1),
        df_transx,
        df_trans1x,
        df_trans2x,
        df_trans3x
    ],
    axis = 1,)

data3x["workclass"] = enc3.fit_transform(golden[["workclass"]])
data3x["marital-status"] = enc3.fit_transform(golden[["marital-status"]])
data3x["education"] = enc3.fit_transform(golden[["education"]])
data3x["sex"] = enc3.fit_transform(golden[["sex"]])

```

```
In [445...] modele = RandomForestClassifier(criterion = 'entropy', max_depth = 5)
```

```
In [446...] modele.fit(data3.drop(['fnlwgt', 'salary'], axis = 1), data3.salary)
```

```
Out[446]: ▼ RandomForestClassifier
RandomForestClassifier(criterion='entropy', max_depth=5)
```

```
In [447...] predictionse = modele.predict(data3x.drop(['fnlwgt', 'salary'], axis = 1))
predictionsex = modele.predict(data3.drop(['fnlwgt', 'salary'], axis = 1))
```

```
In [448...] confusion_matrix(data3.salary, predictionsex)
```

```
Out[448]: array([[23642, 1078],
[ 4005, 3836]], dtype=int64)
```

```
In [449...] print(classification_report(data3.salary, predictionsex))
```

	precision	recall	f1-score	support
<=50K	0.86	0.96	0.90	24720
>50K	0.78	0.49	0.60	7841
accuracy			0.84	32561
macro avg	0.82	0.72	0.75	32561
weighted avg	0.84	0.84	0.83	32561

```
In [460...] ##Occupation
transform_columns4 = ['occupation']
non_num_columns4 = ['relationship', 'race', 'sex', 'native-country']
```

```
In [451...] data4 = df.copy()

onehot4 = preprocessing.OneHotEncoder(handle_unknown="infrequent_if_exist", sparse=False)

enc4 = preprocessing.OrdinalEncoder()

enc4.fit(df[["occupation"]])

transformed4 = onehot4.transform(df[transform_columns4])
new_cols4 = list(onehot4.categories_[0].flatten())
df_trans4 = pd.DataFrame(transformed4, columns = new_cols4)

data4 = pd.concat(
    [
        data4.drop(non_num_columns4, axis = 1),
        df_trans,
        df_trans1,
        df_trans2,
        df_trans3,
        df_trans4
    ],
    axis = 1,)
```

```
data4["workclass"] = enc4.fit_transform(df[["workclass"]])
data4["marital-status"] = enc4.fit_transform(df[["marital-status"]])
data4["education"] = enc4.fit_transform(df[["education"]])
data4["occupation"] = enc4.fit_transform(df[["occupation"]])
data4["sex"] = enc4.fit_transform(df[["sex"]])
```

C:\Users\jorda\anaconda3\lib\site-packages\sklearn\preprocessing_encoders.py:828: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.

warnings.warn(

In [452...

```
data4x = golden.copy()

transformed4 = onehot4.transform(golden[transform_columns4])
new_cols4 = list(onehot4.categories_[0].flatten())
df_trans4x = pd.DataFrame(transformed4, columns = new_cols4)

data4x = pd.concat(
    [
        data4x.drop(non_num_columns4, axis = 1),
        df_transx,
        df_trans1x,
        df_trans2x,
        df_trans3x,
        df_trans4x
    ],
    axis = 1,)

data4x["workclass"] = enc4.fit_transform(golden[["workclass"]])
data4x["marital-status"] = enc4.fit_transform(golden[["marital-status"]])
data4x["education"] = enc4.fit_transform(golden[["education"]])
data4x["occupation"] = enc4.fit_transform(golden[["occupation"]])
data4x["sex"] = enc4.fit_transform(golden[["sex"]])
```

In [453...

```
modelf = RandomForestClassifier(criterion = 'entropy', max_depth = 5)
```

In [454...

```
modelf.fit(data4.drop(['fnlwgt', 'salary'], axis = 1), data4.salary)
```

Out[454]:

```
RandomForestClassifier
RandomForestClassifier(criterion='entropy', max_depth=5)
```

In [455...

```
predictionsf = modelf.predict(data4x.drop(['fnlwgt', 'salary'], axis = 1))
predictionsfx = modelf.predict(data4.drop(['fnlwgt', 'salary'], axis = 1))
```

In [456...

```
confusion_matrix(data4.salary, predictionsfx)
```

Out[456]:

```
array([[24005,  715],
       [ 4329, 3512]], dtype=int64)
```

In [457...

```
print(classification_report(data4.salary, predictionsfx))
```

	precision	recall	f1-score	support
<=50K	0.85	0.97	0.90	24720
>50K	0.83	0.45	0.58	7841
accuracy			0.85	32561
macro avg	0.84	0.71	0.74	32561
weighted avg	0.84	0.85	0.83	32561

In [470...

```
##Relationship
transform_columns5 = ['relationship']
non_num_columns5 = ['race', 'sex', 'native-country']
```

In [461...

```
data5 = df.copy()

onehot5 = preprocessing.OneHotEncoder(handle_unknown="infrequent_if_exist", sparse=False)

enc5 = preprocessing.OrdinalEncoder()

enc5.fit(df[["relationship"]])

transformed5 = onehot5.transform(df[transform_columns5])
new_cols5 = list(onehot5.categories_[0].flatten())
df_trans5 = pd.DataFrame(transformed5, columns = new_cols5)

data5 = pd.concat(
    [
        data5.drop(non_num_columns5, axis = 1),
        df_trans,
        df_trans1,
        df_trans2,
        df_trans3,
        df_trans4,
        df_trans5
    ],
    axis = 1,)

data5["workclass"] = enc5.fit_transform(df[["workclass"]])
data5["marital-status"] = enc5.fit_transform(df[["marital-status"]])
data5["education"] = enc5.fit_transform(df[["education"]])
data5["occupation"] = enc5.fit_transform(df[["occupation"]])
data5["relationship"] = enc5.fit_transform(df[["relationship"]])
data5["sex"] = enc5.fit_transform(df[["sex"]])
```

C:\Users\jorda\anaconda3\lib\site-packages\sklearn\preprocessing_encoders.py:828: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.

warnings.warn(

In [462...

```
data5x = golden.copy()

transformed5 = onehot5.transform(golden[transform_columns5])
new_cols5 = list(onehot5.categories_[0].flatten())
df_trans5x = pd.DataFrame(transformed5, columns = new_cols5)

data5x = pd.concat(
```



```
[
    data5x.drop(non_num_columns5, axis = 1),
    df_transx,
    df_trans1x,
    df_trans2x,
    df_trans3x,
    df_trans4x,
    df_trans5x
],
axis = 1,)

data5x["workclass"] = enc5.fit_transform(golden[["workclass"]])
data5x["marital-status"] = enc5.fit_transform(golden[["marital-status"]])
data5x["education"] = enc5.fit_transform(golden[["education"]])
data5x["occupation"] = enc5.fit_transform(golden[["occupation"]])
data5x["relationship"] = enc5.fit_transform(golden[["relationship"]])
data5x["sex"] = enc5.fit_transform(golden[["sex"]])
```

```
In [463... modelg = RandomForestClassifier(criterion = 'entropy', max_depth = 5)
```

```
In [464... modelg.fit(data5.drop(['fnlwgt', 'salary'], axis = 1), data5.salary)
```

```
Out[464]: ▼ RandomForestClassifier
RandomForestClassifier(criterion='entropy', max_depth=5)
```

```
In [465... predictionsg = modelg.predict(data5x.drop(['fnlwgt', 'salary'], axis = 1))
predictionsgx = modelg.predict(data5.drop(['fnlwgt', 'salary'], axis = 1))
```

```
In [466... confusion_matrix(data5.salary, predictionsgx)
```

```
Out[466]: array([[23697, 1023],
       [ 4016, 3825]], dtype=int64)
```

```
In [467... print(classification_report(data5.salary, predictionsgx))
```

	precision	recall	f1-score	support
<=50K	0.86	0.96	0.90	24720
>50K	0.79	0.49	0.60	7841
accuracy			0.85	32561
macro avg	0.82	0.72	0.75	32561
weighted avg	0.84	0.85	0.83	32561

```
In [481... ##Race
transform_columns6 = ['race']
non_num_columns6 = ['native-country']
```

```
In [486... data6 = df.copy()

onehot6 = preprocessing.OneHotEncoder(handle_unknown="infrequent_if_exist", sparse=False)

enc6 = preprocessing.OrdinalEncoder()

enc6.fit(df[["race"]])
```

```

transformed6 = onehot6.transform(df[transform_columns6])
new_cols6 = list(onehot6.categories_[0].flatten())
df_trans6 = pd.DataFrame(transformed6, columns = new_cols6)

data6 = pd.concat(
    [
        data6.drop(non_num_columns6, axis = 1),
        df_trans,
        df_trans1,
        df_trans2,
        df_trans3,
        df_trans4,
        df_trans5,
        df_trans6
    ],
    axis = 1,)

data6["workclass"] = enc6.fit_transform(df[["workclass"]])
data6["marital-status"] = enc6.fit_transform(df[["marital-status"]])
data6["education"] = enc6.fit_transform(df[["education"]])
data6["occupation"] = enc6.fit_transform(df[["occupation"]])
data6["relationship"] = enc6.fit_transform(df[["relationship"]])
data6["race"] = enc6.fit_transform(df[["race"]])
data6["sex"] = enc6.fit_transform(df[["sex"]])

```

C:\Users\jorda\anaconda3\lib\site-packages\sklearn\preprocessing_encoders.py:828: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.

warnings.warn(

In [487...

```

data6x = golden.copy()

transformed6 = onehot6.transform(golden[transform_columns6])
new_cols6 = list(onehot6.categories_[0].flatten())
df_trans6x = pd.DataFrame(transformed6, columns = new_cols6)

data6x = pd.concat(
    [
        data6x.drop(non_num_columns6, axis = 1),
        df_transx,
        df_trans1x,
        df_trans2x,
        df_trans3x,
        df_trans4x,
        df_trans5x,
        df_trans6x
    ],
    axis = 1,)

data6x["workclass"] = enc6.fit_transform(golden[["workclass"]])
data6x["marital-status"] = enc6.fit_transform(golden[["marital-status"]])
data6x["education"] = enc6.fit_transform(golden[["education"]])
data6x["occupation"] = enc6.fit_transform(golden[["occupation"]])
data6x["relationship"] = enc6.fit_transform(golden[["relationship"]])
data6x["race"] = enc6.fit_transform(golden[["race"]])
data6x["sex"] = enc6.fit_transform(golden[["sex"]])

```

```
In [488... modelh = RandomForestClassifier(criterion = 'entropy', max_depth = 5)
```

```
In [489... modelh.fit(data6.drop(['fnlwgt', 'salary'], axis = 1), data6.salary)
```

```
Out[489]: ▼ RandomForestClassifier
RandomForestClassifier(criterion='entropy', max_depth=5)
```

```
In [490... predictionsh = modelh.predict(data6x.drop(['fnlwgt', 'salary'], axis = 1))
predictionshx = modelh.predict(data6.drop(['fnlwgt', 'salary'], axis = 1))
```

```
In [491... confusion_matrix(data6.salary, predictionshx)
```

```
Out[491]: array([[23668, 1052],
        [ 4027, 3814]], dtype=int64)
```

```
In [492... print(classification_report(data6.salary, predictionshx))
```

	precision	recall	f1-score	support
<=50K	0.85	0.96	0.90	24720
>50K	0.78	0.49	0.60	7841
accuracy			0.84	32561
macro avg	0.82	0.72	0.75	32561
weighted avg	0.84	0.84	0.83	32561

```
In [493... ##Native Country
transform_columns7 = ['native-country']
```

```
In [495... data7 = df.copy()

onehot7 = preprocessing.OneHotEncoder(handle_unknown="infrequent_if_exist", sparse=False)

enc7 = preprocessing.OrdinalEncoder()

enc7.fit(df[["native-country"]])

transformed7 = onehot7.transform(df[transform_columns7])
new_cols7 = list(onehot7.categories_[0].flatten())
df_trans7 = pd.DataFrame(transformed7, columns = new_cols7)

data7 = pd.concat(
    [
        data7,
        df_trans,
        df_trans1,
        df_trans2,
        df_trans3,
        df_trans4,
        df_trans5,
        df_trans6,
        df_trans7
    ],
```

```
axis = 1,)
```

```
data7["workclass"] = enc7.fit_transform(df[["workclass"]])
data7["marital-status"] = enc7.fit_transform(df[["marital-status"]])
data7["education"] = enc7.fit_transform(df[["education"]])
data7["occupation"] = enc7.fit_transform(df[["occupation"]])
data7["relationship"] = enc7.fit_transform(df[["relationship"]])
data7["race"] = enc7.fit_transform(df[["race"]])
data7["sex"] = enc7.fit_transform(df[["sex"]])
data7["native-country"] = enc7.fit_transform(df[["native-country"]])
```

C:\Users\jorda\anaconda3\lib\site-packages\sklearn\preprocessing_encoders.py:828: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.

```
warnings.warn(
```

In [496...

```
data7x = golden.copy()
```

```
transformed7 = onehot7.transform(golden[transform_columns7])
new_cols7 = list(onehot7.categories_[0].flatten())
df_trans7x = pd.DataFrame(transformed7, columns = new_cols7)
```

```
data7x = pd.concat(
    [
        data7x,
        df_transx,
        df_trans1x,
        df_trans2x,
        df_trans3x,
        df_trans4x,
        df_trans5x,
        df_trans6x,
        df_trans7x
    ],
    axis = 1,)
```

```
data7x["workclass"] = enc7.fit_transform(golden[["workclass"]])
data7x["marital-status"] = enc7.fit_transform(golden[["marital-status"]])
data7x["education"] = enc7.fit_transform(golden[["education"]])
data7x["occupation"] = enc7.fit_transform(golden[["occupation"]])
data7x["relationship"] = enc7.fit_transform(golden[["relationship"]])
data7x["race"] = enc7.fit_transform(golden[["race"]])
data7x["sex"] = enc7.fit_transform(golden[["sex"]])
data7x["native-country"] = enc7.fit_transform(golden[["native-country"]])
```

In [497...

```
model1 = RandomForestClassifier(criterion = 'entropy', max_depth = 5)
```

In [498...

```
model1.fit(data7.drop(['fnlwgt', 'salary'], axis = 1), data7.salary)
```

Out[498]:

```
RandomForestClassifier
RandomForestClassifier(criterion='entropy', max_depth=5)
```

In [500...

```
predictionsi = model1.predict(data7x.drop(['fnlwgt', 'salary'], axis = 1))
predictionsix = model1.predict(data7.drop(['fnlwgt', 'salary'], axis = 1))
```

In [501...

```
confusion_matrix(data7.salary, predictionsix)
```

```
Out[501]: array([[23641, 1079],  
                [ 4210, 3631]], dtype=int64)
```

```
In [502... print(classification_report(data7.salary, predictionsix))
```

	precision	recall	f1-score	support
<=50K	0.85	0.96	0.90	24720
>50K	0.77	0.46	0.58	7841
accuracy			0.84	32561
macro avg	0.81	0.71	0.74	32561
weighted avg	0.83	0.84	0.82	32561

```
In [ ]:
```